



Binary Classification with a Bank Churn – Competition Notebook (Prediction Submission)

Sawandi Kirby
2024-01-25

Contents

1	Ask 5	
1.1	About the Data	5
2	Prepare	5
2.1	Get Enviornment ready	5
2.2	Trim rows	9
2.3	Check for missing values	9
3	Analyze	10
3.1	Train Data	10
3.2	Train logistic regression model	13
3.3	Train XGBoost model	14

3.4	Create matrix with features and labels	15
4	Predictions	17
4.1	XGBoost	17
4.2	Logistical Model	18
5	Share	19
5.1	Visual	19
5.2	Archive	20
6	Logistic Regrsson Submission Entry Score = 0.80891	20
7	XGBoost Submission Entry Score = .88418	21

Ask

About the Data

The bank customer churn dataset is a commonly used dataset for predicting customer churn in the banking industry. It contains information on bank customers who either left the bank or continue to be a customer.

Our task is to predict whether a customer continues with their account or closes it (e.g., churns).

Prepare

Get Environment ready

Load required libraries

```
library(tidyverse)

## — Attaching core tidyverse packages
tidyverse 2.0.0 —
## ✓ dplyr      1.1.4      ✓ readr      2.1.5
## ✓ forcats   1.0.0      ✓ stringr    1.5.1
## ✓ ggplot2    3.5.1      ✓ tibble     3.2.1
## ✓ lubridate  1.9.3      ✓ tidyr      1.3.1
## ✓ purrr      1.0.2
## — Conflicts

tidyverse_conflicts() —
## ✗ dplyr::filter() masks stats::filter()
## ✗ dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all
conflicts to become errors

library(ggplot2)
library(corrplot)

## corrplot 0.92 loaded

library(caret)

## Loading required package: lattice
##
## Attaching package: 'caret'
##
## The following object is masked from 'package:purrr':
##
##   lift
```

```
library(glmnet)

## Loading required package: Matrix
##
## Attaching package: 'Matrix'
##
## The following objects are masked from 'package:tidyr':
##
##     expand, pack, unpack
##
## Loaded glmnet 4.1-8
```

```
library(pROC)

## Type 'citation("pROC")' for a citation.
##
## Attaching package: 'pROC'
##
## The following objects are masked from 'package:stats':
##
##     cov, smooth, var
```

```
library(xgboost)

##
## Attaching package: 'xgboost'
##
## The following object is masked from 'package:dplyr':
##
##     slice
```

Create function to apply trimws to character columns of a data frame

```
trimws_df <- function(df) {
  char_cols <- sapply(df, is.character)
  df[char_cols] <- lapply(df[char_cols], trimws)
  return(df)
}
```

Load data

```
test <- read_csv("playground-series-s4e1\\test.csv")
```

```
## Rows: 110023 Columns: 13
## — Column specification
```

```
## Delimiter: ","
## chr (3): Surname, Geography, Gender
## dbl (10): id, CustomerId, CreditScore, Age, Tenure, Balance,
## NumOfProducts, HasCrCa...
##
```

```
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this
message.
```

```
train <- read_csv("playground-series-s4e1\\train.csv")
```

```
## Rows: 165034 Columns: 14
## — Column specification
```

```
## Delimiter: ","
## chr (3): Surname, Geography, Gender
## dbl (11): id, CustomerId, CreditScore, Age, Tenure, Balance,
NumOfProducts, HasCrCa...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this
message.
```

```
sample_submission <-
read_csv("playground-series-s4e1\\sample_submission.csv")
```

```
## Rows: 110023 Columns: 2
## — Column specification
```

```
## Delimiter: ","
## dbl (2): id, Exited
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this
message.
```

Exploratory Data Analysis (EDA)

Get a glimpse of the datasets

```
head(train)
```

```
## # A tibble: 6 × 14
##       id CustomerId Surname      CreditScore Geography Gender    Age
Tenure Balance
##   <dbl>      <dbl> <chr>          <dbl> <chr>      <chr> <dbl>
<dbl>  <dbl>
## 1      0   15674932 Okwudilichukwu      668 France    Male     33
3      0
## 2      1   15749177 Okwudiliolisa      627 France    Male     33
1      0
## 3      2   15694510 Hsueh              678 France    Male     40
10     0
## 4      3   15741417 Kao                581 France    Male     34
2 148883.
```

```
## 5      4  15766172 Chiemenam          716 Spain      Male      33
5      0
## 6      5  15771669 Genovese          588 Germany    Male      36
4 131779.
## # i 5 more variables: NumOfProducts <dbl>, HasCrCard <dbl>, IsActiveMember
<dbl>,
## #   EstimatedSalary <dbl>, Exited <dbl>
```

`head(test)`

```
## # A tibble: 6 × 13
##       id CustomerId Surname   CreditScore Geography Gender   Age Tenure
Balance
##   <dbl>      <dbl> <chr>         <dbl> <chr>      <chr> <dbl> <dbl>
<dbl>
## 1 165034   15773898 Lucchese         586 France    Female   23     2
0
## 2 165035   15782418 Nott             683 France    Female   46     2
0
## 3 165036   15807120 K?             656 France    Female   34     7
0
## 4 165037   15808905 O'Donnell       681 France    Male     36     8
0
## 5 165038   15607314 Higgins         752 Germany   Male     38    10
121264.
## 6 165039   15672704 Pearson         593 France    Female   22     9
0
## # i 4 more variables: NumOfProducts <dbl>, HasCrCard <dbl>, IsActiveMember
<dbl>,
## #   EstimatedSalary <dbl>
```

`head(sample_submission)`

```
## # A tibble: 6 × 2
##       id Exited
##   <dbl> <dbl>
## 1 165034    0.5
## 2 165035    0.5
## 3 165036    0.5
## 4 165037    0.5
## 5 165038    0.5
## 6 165039    0.5
```

Review the structure of the train data

`str(train)`

```
## spc_tbl_ [165,034 × 14] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ id      : num [1:165034] 0 1 2 3 4 5 6 7 8 9 ...
## $ CustomerId : num [1:165034] 15674932 15749177 15694510 15741417
15766172 ...
## $ Surname   : chr [1:165034] "Okwudilichukwu" "Okwudiliolisa"
```

```

"Hsueh" "Kao" ...
## $ CreditScore : num [1:165034] 668 627 678 581 716 588 593 678 676 583
...
## $ Geography : chr [1:165034] "France" "France" "France" "France" ...
## $ Gender : chr [1:165034] "Male" "Male" "Male" "Male" ...
## $ Age : num [1:165034] 33 33 40 34 33 36 30 37 43 40 ...
## $ Tenure : num [1:165034] 3 1 10 2 5 4 8 1 4 4 ...
## $ Balance : num [1:165034] 0 0 0 148883 0 ...
## $ NumOfProducts : num [1:165034] 2 2 2 1 2 1 1 1 2 1 ...
## $ HasCrCard : num [1:165034] 1 1 1 1 1 1 1 1 1 1 ...
## $ IsActiveMember : num [1:165034] 0 1 0 1 1 0 0 0 0 1 ...
## $ EstimatedSalary: num [1:165034] 181450 49504 184867 84561 15069 ...
## $ Exited : num [1:165034] 0 0 0 0 0 1 0 0 0 0 ...
## - attr(*, "spec")=
## .. cols(
## .. id = col_double(),
## .. CustomerId = col_double(),
## .. Surname = col_character(),
## .. CreditScore = col_double(),
## .. Geography = col_character(),
## .. Gender = col_character(),
## .. Age = col_double(),
## .. Tenure = col_double(),
## .. Balance = col_double(),
## .. NumOfProducts = col_double(),
## .. HasCrCard = col_double(),
## .. IsActiveMember = col_double(),
## .. EstimatedSalary = col_double(),
## .. Exited = col_double()
## .. )
## - attr(*, "problems")=<externalptr>

```

The dataset includes the following attributes:

Customer ID: A unique identifier for each customer

Surname: The customer's surname or last name

Credit Score: A numerical value representing the customer's credit score

Geography: The country where the customer resides (France, Spain or Germany)

Gender: The customer's gender (Male or Female)

Age: The customer's age.

Tenure: The number of years the customer has been with the bank

Balance: The customer's account balance

NumOfProducts: The number of bank products the customer uses (e.g., savings ### account, credit card)

HasCrCard: Whether the customer has a credit card (1 = yes, 0 = no)

IsActiveMember: Whether the customer is an active member (1 = yes, 0 = no)

EstimatedSalary: The estimated salary of the customer

Exited: Whether the customer has churned (1 = yes, 0 = no)

Trim rows

Use trimrows function created previously

This will go through the data frame find the character columns and trim whitespaces around the values

```
train <- trimws_df(train)
```

Check for missing values

```
missing_values <- colSums(is.na(train))  
head(missing_values)
```

```
##          id CustomerId  Surname CreditScore  Geography  Gender  
##          0           0         0           0           0         0
```


Analyze

Train Data

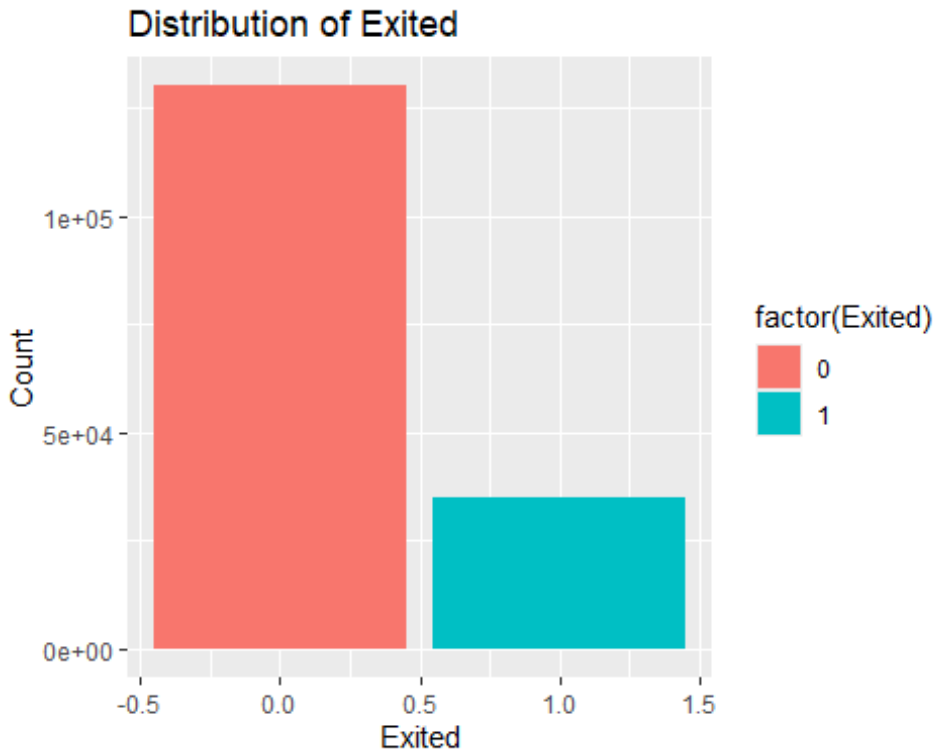
View summary statistics of train data

```
summary(train)
```

```
##      id      CustomerId      Surname      CreditScore
## Min.   :    0      Min.   :15565701      Length:165034      Min.   :350.0
## 1st Qu.: 41258      1st Qu.:15633141      Class :character      1st Qu.:597.0
## Median : 82517      Median :15690169      Mode  :character      Median :659.0
## Mean   : 82517      Mean   :15692005                      Mean   :656.5
## 3rd Qu.:123775      3rd Qu.:15756824                      3rd Qu.:710.0
## Max.   :165033      Max.   :15815690                      Max.   :850.0
## Geography      Gender      Age      Tenure
## Length:165034      Length:165034      Min.   :18.00      Min.   : 0.00
## Class :character      Class :character      1st Qu.:32.00      1st Qu.: 3.00
## Mode  :character      Mode  :character      Median :37.00      Median : 5.00
##                      Mean   :38.13      Mean   : 5.02
##                      3rd Qu.:42.00      3rd Qu.: 7.00
##                      Max.   :92.00      Max.   :10.00
##      Balance      NumOfProducts      HasCrCard      IsActiveMember
## EstimatedSalary
## Min.   :    0      Min.   :1.000      Min.   :0.000      Min.   :0.0000      Min.
## :    11.58
## 1st Qu.:    0      1st Qu.:1.000      1st Qu.:1.000      1st Qu.:0.0000      1st
## Qu.: 74637.57
## Median :    0      Median :2.000      Median :1.000      Median :0.0000      Median
## :117948.00
## Mean   : 55478      Mean   :1.554      Mean   :0.754      Mean   :0.4978      Mean
## :112574.82
## 3rd Qu.:119940      3rd Qu.:2.000      3rd Qu.:1.000      3rd Qu.:1.0000      3rd
## Qu.:155152.47
## Max.   :250898      Max.   :4.000      Max.   :1.000      Max.   :1.0000      Max.
## :199992.48
##      Exited
## Min.   :0.0000
## 1st Qu.:0.0000
## Median :0.0000
## Mean   :0.2116
## 3rd Qu.:0.0000
## Max.   :1.0000
```

Visualize relationships using ggplot2

```
ggplot(train, aes(x = Exited, fill = factor(Exited))) +
  geom_bar() +
  labs(title = "Distribution of Exited",
       x = "Exited",
       y = "Count")
```



Explore relationships with 'Exited' for numerical variables


```
cor(train[, c("CreditScore", "Age", "Tenure", "Balance", "NumOfProducts",
              "HasCrCard", "IsActiveMember", "EstimatedSalary", "Exited")])
```

##	CreditScore	Age	Tenure	Balance
NumOfProducts				
## CreditScore	1.0000000000	-0.008918146	0.0009424799	0.006973053
0.011360808				
## Age	-0.0089181458	1.0000000000	-0.0108303451	0.064318287
-0.102194912				
## Tenure	0.0009424799	-0.010830345	1.0000000000	-0.009481186
0.007334828				
## Balance	0.0069730535	0.064318287	-0.0094811862	1.0000000000
-0.361032521				
## NumOfProducts	0.0113608082	-0.102194912	0.0073348275	-0.361032521
1.0000000000				
## HasCrCard	-0.0028277567	-0.012111332	0.0053266159	-0.018584007
0.005482281				
## IsActiveMember	0.0147902638	0.003319563	-0.0055322590	-0.015073487
0.039736070				
## EstimatedSalary	-0.0018203035	-0.005398663	0.0009705869	0.008586201
-0.004285089				
## Exited	-0.0273826001	0.340768163	-0.0195648445	0.129742860
-0.214554232				
##	HasCrCard	IsActiveMember	EstimatedSalary	Exited
## CreditScore	-0.002827757	0.014790264	-0.0018203035	-0.02738260

```
## Age          -0.012111332    0.003319563   -0.0053986634   0.34076816
## Tenure       0.005326616   -0.005532259    0.0009705869  -0.01956484
## Balance     -0.018584007   -0.015073487    0.0085862012   0.12974286
## NumOfProducts 0.005482281    0.039736070   -0.0042850891  -0.21455423
## HasCrCard    1.000000000   -0.021033789    0.0044382187  -0.02214133
## IsActiveMember -0.021033789    1.000000000   -0.0080800461  -0.21023703
## EstimatedSalary 0.004438219   -0.008080046    1.0000000000   0.01882681
## Exited      -0.022141333   -0.210237026    0.0188268057   1.00000000
```

Build a heatmap for better visualization of correlations

```
cor_data <- train[, c("CreditScore", "Age", "Tenure", "Balance",
"NumOfProducts", "HasCrCard", "IsActiveMember", "EstimatedSalary", "Exited")]
cor_matrix <- cor(cor_data)
corrplot(cor_matrix, method = "color", type = "upper", order = "hclust",
tl.cex = 0.7)
```



View character categorical variables

```
table(train$Geography, train$Exited)
```

```
##
##           0      1
##  France 78643 15572
##  Germany 21492 13114
##   Spain  29978  6235
```

```
table(train$Gender, train$Exited)
```

```
##
##           0      1
## Female 51779 20105
## Male   78334 14816
```

Convert character categorical variables to factors

```
train_encoded <- train %>%
  select(-c(id, CustomerId, Surname)) %>%
  mutate(Geography = as.factor(Geography),
         Gender = as.factor(Gender),
         Exited = as.factor(Exited))
```

Split data into features (X) and target variable (y)

```
X <- select(train_encoded, -Exited)
y <- train_encoded$Exited
```

Split data into training and testing sets

```
set.seed(123)
split_index <- createDataPartition(y, p = 0.7, list = FALSE)
train_data <- train_encoded[split_index, ]
test_data <- train_encoded[-split_index, ]
```

```
Log_train_data <- train_encoded[split_index, ]
Log_test_data <- train_encoded[-split_index, ]
```

```
XGB_train_data <- train_encoded[split_index, ]
XGB_test_data <- train_encoded[-split_index, ]
```

Turn Exited column into factor

```
Log_train_data$Exited <- as.factor(Log_train_data$Exited)
XGB_train_data$Exited <- as.factor(XGB_train_data$Exited)
```

Train logistic regression model

```
logistic_model <- glm(Exited ~ ., data = Log_train_data, family = "binomial")
summary(logistic_model)
```

```
##
## Call:
## glm(formula = Exited ~ ., family = "binomial", data = Log_train_data)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -2.430e+00  8.784e-02 -27.662 < 2e-16 ***
## CreditScore  -7.890e-04  1.039e-04  -7.596 3.04e-14 ***
## GeographyGermany  1.154e+00  2.360e-02  48.882 < 2e-16 ***
## GeographySpain    4.251e-02  2.198e-02   1.934  0.0531 .
## GenderMale      -6.822e-01  1.675e-02 -40.723 < 2e-16 ***
## Age             9.430e-02  9.440e-04  99.894 < 2e-16 ***
## Tenure          -1.794e-02  2.965e-03  -6.051 1.44e-09 ***
```

```
## Balance          -1.971e-06  1.701e-07 -11.583  < 2e-16 ***
## NumOfProducts    -9.154e-01  1.650e-02 -55.490  < 2e-16 ***
## HasCrCard        -1.549e-01  1.914e-02  -8.093  5.83e-16 ***
## IsActiveMember   -1.293e+00  1.799e-02 -71.871  < 2e-16 ***
## EstimatedSalary   9.597e-07  1.663e-07   5.771  7.87e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 119237  on 115524  degrees of freedom
## Residual deviance:  91122  on 115513  degrees of freedom
## AIC: 91146
##
## Number of Fisher Scoring iterations: 5
```

Predictions on train_data using logistic regression

```
predictions_train_logistic <- predict(logistic_model, newdata =
  Log_train_data, type = "response")
binary_predictions_train_logistic <- ifelse(predictions_train_logistic > 0.5,
  1, 0)
train_data$PredictedExited_Logistic <- binary_predictions_train_logistic
```

Train XGBoost model

```
xgb_features <- setdiff(names(XGB_train_data), c("id", "Surname", "Exited",
  "PredictedExited_Logistic"))
```

Ensure relevant columns are numeric

```
XGB_train_data[, xgb_features] <- lapply(XGB_train_data[, xgb_features],
  as.numeric)
```

Convert Exited to integer

```
XGB_train_data$Exited <- as.integer(XGB_train_data$Exited) - 1
```

Create matrix with features and labels

```
xgb_matrix <- xgb.DMatrix(as.matrix(XGB_train_data[, xgb_features]), label =
  XGB_train_data$Exited)
```

Create xgboost model

```
xgb_model <- xgboost(
  data = xgb_matrix,
  nrounds = 100,
  objective = "binary:logistic",
  eval_metric = "logloss",
  verbose = 1
)
```

```
## [1] train-logloss:0.543701
## [2] train-logloss:0.462336
## [3] train-logloss:0.412461
## [4] train-logloss:0.380961
## [5] train-logloss:0.359584
## [6] train-logloss:0.345226
## [7] train-logloss:0.335974
## [8] train-logloss:0.329861
## [9] train-logloss:0.325658
## [10] train-logloss:0.322307
## [11] train-logloss:0.320131
## [12] train-logloss:0.318390
## [13] train-logloss:0.316737
## [14] train-logloss:0.315470
## [15] train-logloss:0.314520
## [16] train-logloss:0.313604
## [17] train-logloss:0.312482
## [18] train-logloss:0.311714
## [19] train-logloss:0.311183
## [20] train-logloss:0.310581
## [21] train-logloss:0.309813
## [22] train-logloss:0.309084
## [23] train-logloss:0.308712
## [24] train-logloss:0.308201
## [25] train-logloss:0.307895
## [26] train-logloss:0.307645
## [27] train-logloss:0.307268
## [28] train-logloss:0.306740
## [29] train-logloss:0.306295
## [30] train-logloss:0.305839
## [31] train-logloss:0.305605
## [32] train-logloss:0.305331
## [33] train-logloss:0.304996
## [34] train-logloss:0.304396
## [35] train-logloss:0.303979
## [36] train-logloss:0.303414
## [37] train-logloss:0.303012
## [38] train-logloss:0.302704
## [39] train-logloss:0.302152
## [40] train-logloss:0.301955
## [41] train-logloss:0.301361
## [42] train-logloss:0.301265
## [43] train-logloss:0.300832
## [44] train-logloss:0.300524
## [45] train-logloss:0.300114
## [46] train-logloss:0.299660
## [47] train-logloss:0.299616
## [48] train-logloss:0.299285
## [49] train-logloss:0.298938
## [50] train-logloss:0.298681
```

```
## [51] train-logloss:0.298142
## [52] train-logloss:0.298068
## [53] train-logloss:0.297926
## [54] train-logloss:0.297558
## [55] train-logloss:0.297218
## [56] train-logloss:0.297096
## [57] train-logloss:0.296917
## [58] train-logloss:0.296738
## [59] train-logloss:0.296551
## [60] train-logloss:0.296179
## [61] train-logloss:0.295673
## [62] train-logloss:0.295116
## [63] train-logloss:0.294991
## [64] train-logloss:0.294726
## [65] train-logloss:0.294525
## [66] train-logloss:0.294153
## [67] train-logloss:0.293980
## [68] train-logloss:0.293652
## [69] train-logloss:0.293333
## [70] train-logloss:0.293068
## [71] train-logloss:0.292707
## [72] train-logloss:0.292461
## [73] train-logloss:0.292430
## [74] train-logloss:0.292045
## [75] train-logloss:0.291600
## [76] train-logloss:0.291247
## [77] train-logloss:0.291174
## [78] train-logloss:0.290648
## [79] train-logloss:0.290502
## [80] train-logloss:0.290215
## [81] train-logloss:0.289993
## [82] train-logloss:0.289767
## [83] train-logloss:0.289315
## [84] train-logloss:0.289098
## [85] train-logloss:0.288972
## [86] train-logloss:0.288591
## [87] train-logloss:0.288351
## [88] train-logloss:0.287893
## [89] train-logloss:0.287826
## [90] train-logloss:0.287453
## [91] train-logloss:0.286964
## [92] train-logloss:0.286841
## [93] train-logloss:0.286783
## [94] train-logloss:0.286579
## [95] train-logloss:0.286352
## [96] train-logloss:0.286126
## [97] train-logloss:0.286106
## [98] train-logloss:0.285915
## [99] train-logloss:0.285555
## [100] train-logloss:0.285188
```

Make predictions using XGBoost model

Convert probabilities to binary predictions

```
xgb_predictions_train <- predict(xgb_model, newdata = xgb_matrix)
binary_xgb_predictions_train <- ifelse(xgb_predictions_train > 0.5, 1, 0)
train_data$PredictedExited_XGBoost <- binary_xgb_predictions_train
```

Compare the 'Exited' columns for accuracy

```
accuracy_logistic <- sum(train_data$Exited ==
  train_data$PredictedExited_Logistic) / nrow(train_data)
accuracy_xgb <- sum(train_data$Exited == binary_xgb_predictions_train) /
  nrow(train_data)

cat("Logistic Regression Accuracy on train_data:", round(accuracy_logistic *
  100, 2), "%\n")

## Logistic Regression Accuracy on train_data: 83.44 %

cat("XGBoost Accuracy on train_data:", round(accuracy_xgb * 100, 2), "%\n")

## XGBoost Accuracy on train_data: 88.07 %
```

Predictions

XGBoost

Convert character categorical variables to factors in test data

```
test_XGB <- test %>%
  select(-c(id, CustomerId, Surname)) %>%
  mutate(Geography = as.factor(Geography),
    Gender = as.factor(Gender))
```

Ensure relevant columns are numeric

```
test_XGB[, xgb_features] <- lapply(test_XGB[, xgb_features], as.numeric)
```

Create a matrix for our predictions

Run the XGBoost prediction model on the test_XGB data set

```
xgb_matrix_test <- xgb.DMatrix(as.matrix(test_XGB[, xgb_features]))
xgb_predictions_test <- predict(xgb_model, newdata = xgb_matrix_test)
binary_xgb_predictions_test <- ifelse(xgb_predictions_test > 0.5, 1, 0)
test_XGB$PredictedExited_XGBoost <- binary_xgb_predictions_test
```

Evaluate the model on test data using ROC curve

```
roc_curve_test <- roc(test_XGB$PredictedExited_XGBoost, xgb_predictions_test)

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases
```


Logistical Model

Create a new data frame for our logistical model predictions

```
test_Logistic <- test %>%  
  select(-c(id, CustomerId, Surname)) %>%  
  mutate(Geography = as.factor(Geography),  
         Gender = as.factor(Gender))
```

Run the logistical predictions on the test_Logistic data set

```
predictions_test_logistic <- predict(logistic_model, newdata = test_Logistic,  
type = "response")  
binary_predictions_test_logistic <- ifelse(predictions_test_logistic > 0.5,  
1, 0)  
test_Logistic$PredictedExited_Logistic <- binary_predictions_test_logistic
```

Evaluate the model on test data using ROC curve for logistic regression

```
roc_curve_logistic_test <- roc(test_Logistic$PredictedExited_Logistic,  
predictions_test_logistic)  
  
## Setting levels: control = 0, case = 1  
  
## Setting direction: controls < cases
```

Share

Visual

Summary of ROC curves

```
summary(roc_curve_test)
```

```
##               Length Class  Mode  
## percent              1 -none- logical  
## sensitivities       108248 -none- numeric  
## specificities       108248 -none- numeric  
## thresholds          108248 -none- numeric  
## direction              1 -none- character  
## cases                 17466 -none- numeric  
## controls              92557 -none- numeric  
## fun.sesp              1 -none- function  
## auc                   1 auc    numeric  
## call                  3 -none- call  
## original.predictor  110023 -none- numeric  
## original.response  110023 -none- numeric  
## predictor           110023 -none- numeric  
## response            110023 -none- numeric  
## levels               2 -none- character
```

```
summary(roc_curve_logistic_test)
```

```
##           Length Class  Mode
## percent           1 -none- logical
## sensitivities    109952 -none- numeric
## specificities    109952 -none- numeric
## thresholds       109952 -none- numeric
## direction         1 -none- character
## cases            12970 -none- numeric
## controls          97053 -none- numeric
## fun.sesp           1 -none- function
## auc                1 auc    numeric
## call              3 -none- call
## original.predictor 110023 -none- numeric
## original.response  110023 -none- numeric
## predictor         110023 -none- numeric
## response          110023 -none- numeric
## levels            2 -none- character
```

AUC of ROC curves

```
auc_test <- auc(roc_curve_test)
auc_logistic_test <- auc(roc_curve_logistic_test)
```

```
cat("XGBoost AUC on test data:", auc_test, "\n")
```

```
## XGBoost AUC on test data: 1
```

```
cat("Logistic Regression AUC on test data:", auc_logistic_test, "\n")
```

```
## Logistic Regression AUC on test data: 1
```

Archive

Create submission data frames

```
Logistic_submission <- data.frame(id = test$id, Exited =
  predictions_test_logistic)
```

```
XGBoost_submission <- data.frame(id = test$id, Exited = xgb_predictions_test)
```


Logistic Regression Submission Entry Score = 0.80891

Review the submission data set

```
head(Logistic_submission)
```

```
##      id      Exited
## 1 165034 0.02346466
## 2 165035 0.58237170
## 3 165036 0.15197213
## 4 165037 0.20440605
```

```
## 5 165038 0.41989777
## 6 165039 0.05921289
```

 SAWANDI KIRBY · 4MO AGO · 855 VIEWS


▲ 8

Edit

(R) Binary Classification with a Bank Churn

R · [submission](#), [Binary Classification with a Bank Churn Dataset](#)

Notebook Input Output Logs Comments (0) Settings

 4.01

Competition Notebook

Run

Private Score

Public Score

Best Score

Version 3 of 3

[Binary Classification with a Bank Churn ...](#)23.8s0.813270.808910.81327 V3


Logistic Submission

XGBoost Submission Entry Score = .88418

Review the submission data set

`head(XGBoost_submission)`

```
##      id      Exited
## 1 165034 0.04157397
## 2 165035 0.84078336
## 3 165036 0.02030335
## 4 165037 0.23674323
## 5 165038 0.37523997
## 6 165039 0.03150999
```

 SAWANDI KIRBY · 4MO AGO · 97 VIEWS


▲ 8

Edit

(R) XGBoost Bank Churn Binary Classification

R · [XGB submission unrrounded](#), [Binary Classification with a Bank Churn Dataset](#)

Notebook Input Output Logs Comments (2) Settings

 4.01

Competition Notebook

Run

Private Score

Public Score

Best Score

Version 4 of 4

[Binary Classification with a Bank Churn ...](#)31.7s0.885170.884180.88517 V4

XGBoost Submission